

Fundamentals of Data Engineering

Week 01 - sync session

datascience@berkeley

Week 1 - Overview

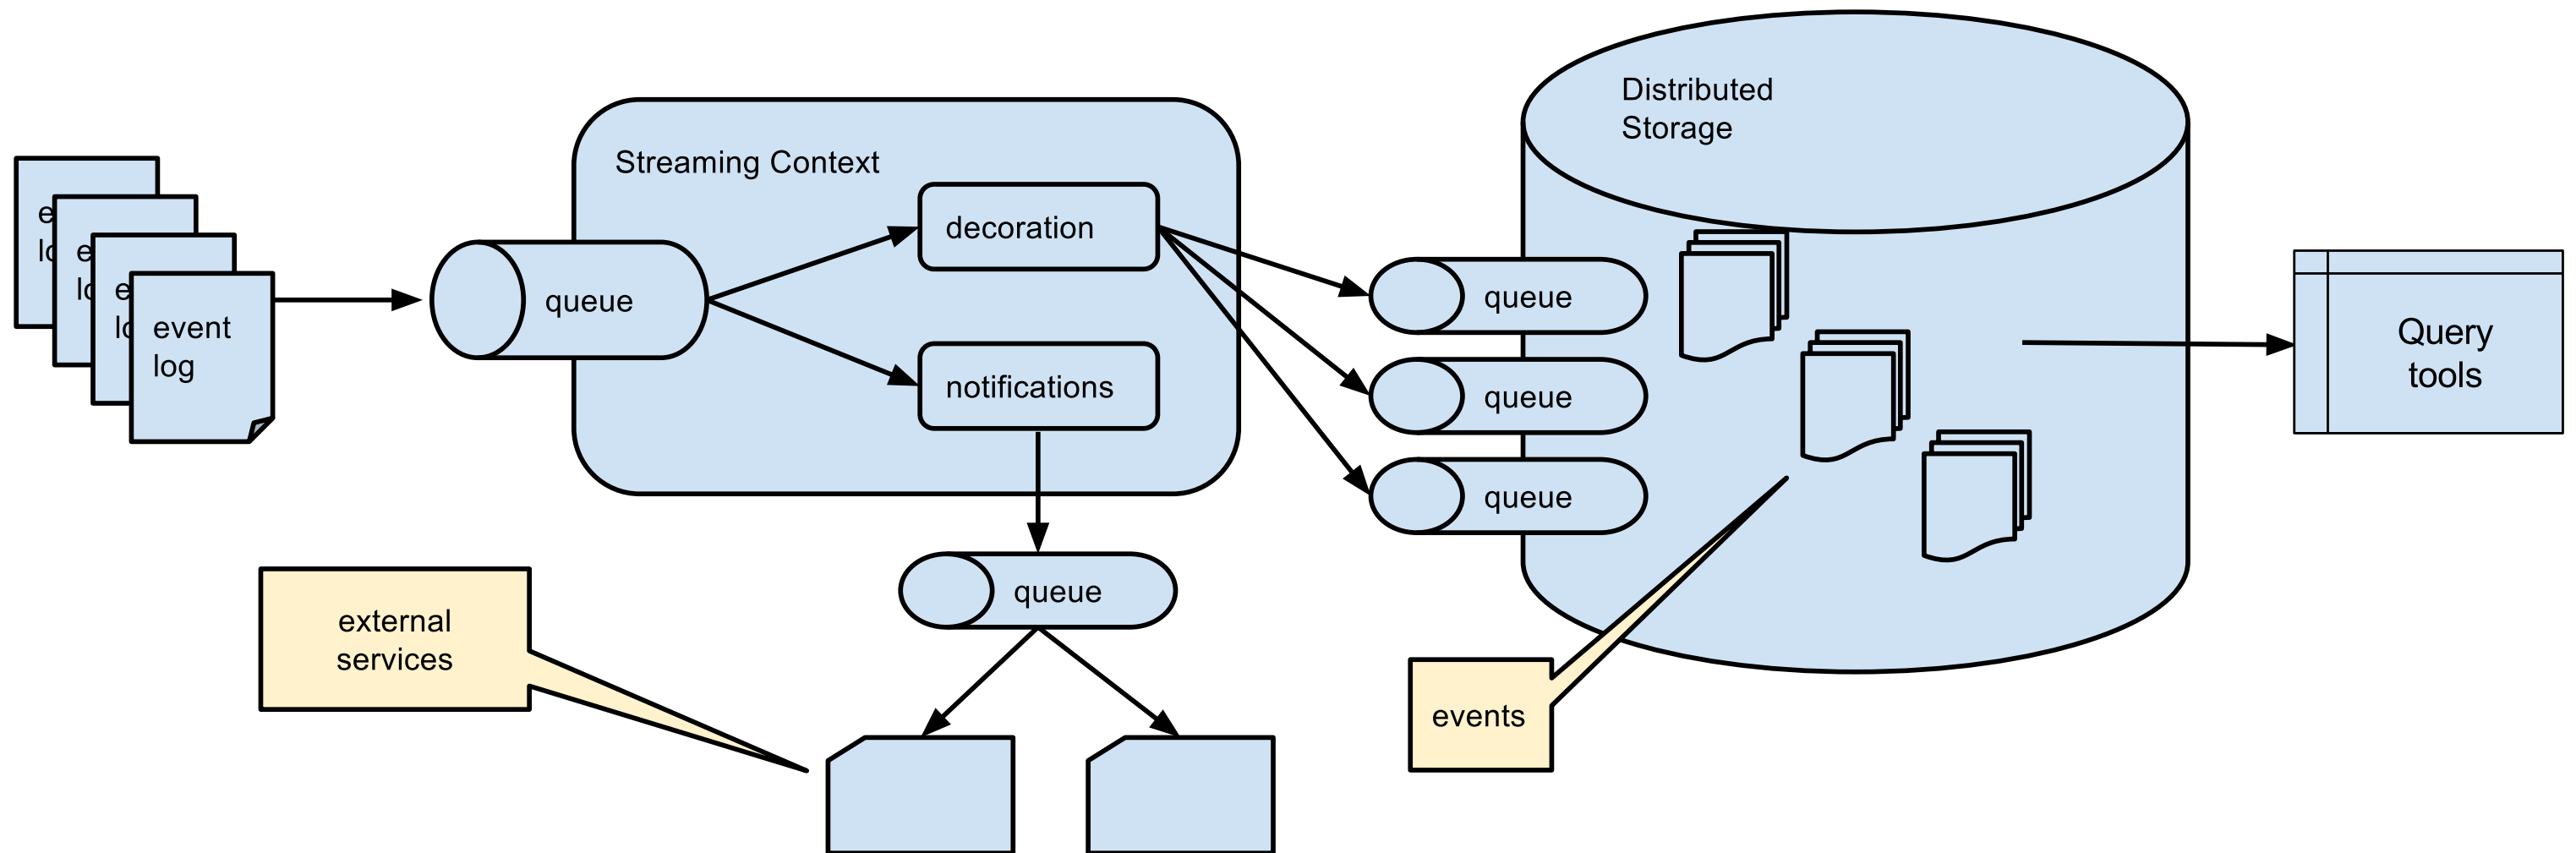
- Introductions
- Set up your working environment for this class
- Review syllabus, course goals, processes & tools ...

Introductions

In this class, you will

- Gain exposure to basic problems associated with data and data-driven decision-making
- Develop a working knowledge of some tools/techniques used to solve these problems
- Learn where to go for help and more info

Just enough



Process/Procedures

- Good practices
- Appropriate tools
- Getting used to

Activities

- Let's get going!

Slack

Cloud Instances

Access

- URL:

```
http://<ip_address>:8000/
```

- Login: science
- Password: (provided over slack)

JupyterHub Terminal

Docker



- pull the image:

```
docker pull midsw205/base
```

- create your midsw205 workspace:

```
mkdir w205
```

- run (set *your* home directory for “-v”)

```
docker run \  
  -it \  
  --rm \  
  -v /home/science/w205:/w205 \  
  midsw205/base:latest \  
  bash
```

- exit (or ctrl-d)

Git

Git set up

Clone the course-content repo

- `cd w205`
- Clone the course-content repo into your midsw205 workspace:

```
git clone https://github.com/mids-w205- \  
  <instructor-last-name>/course-content
```

Signup Assignment

Clone the repo

- `cd w205`
- Clone the repo into your mids-w205 workspace:

```
git clone https://github.com/mids-w205- \  
  <instructor-last-name>/ \  
  signup-<git-user-name>
```

Open, Change, Close README.md

- `nano README.md`
- change line
- `ctrl-o`
- return
- `ctrl-x`
- Now you're out of nano.

Git: commit changes

- `git status`
- `git add README.md`
- `git commit -m 'my new readme'`
- The first time you commit, it doesn't know who you are.

```
git config --global user.email "you@example.com"
```

```
git config --global user.name "Your Name"
```

- `git commit -m 'my new readme'`
- `git push`

Git: submit a PR

- All assignments submitted as PRs

```
https://github.com/mids-w205-martin-mims/signup-<user-name>
```

- Click on README .md
- Click on edit button (pencil icon)
- Make a change
- “Commit changes” section, select “Create a new branch for this commit...”
- Enter PR name & description
- Click “Propose file change” button
- Assign instructors as reviewers
- Click “Create pull request” button

How this class works

Syllabus

<https://github.com/mids-w205-martin-mims/course-content>

Asynchronous Content

```
https://github.com/mids-w205-martin-mims/course-content/ \
blob/master/01-Introduction/async-videos.md
```

- Same as in ISVC, but you can access it all in one place here.

Readings

- No one textbook available for this course.
- Using subscription service to cover the range of topics.
- <https://www.safaribooksonline.com/pricing/>
- Individual option: \$39/month (can stop whenever you want)
- Quick note: Get the mobile apps.

Prerequisites

- Resources listed under prereqs
- Safari has tons of other materials you can help yourself with.

Course Outline

- 4 sections:
- 3-week Introduction
- 5-week Basics section
- 4-week Streaming Data section
- Putting it All Together

Class flow

Class 1

- Preview, discussion, walkthrough set up for github for Assignment 1

Between Class 1 & Class 2

- async material for Week 1
- Readings for Week 1
- Turn in Assignment 01

Class 2

- Groups share Assignment 01
- Preview Query Project
- Preview, discussion, SQL query activities to prepare for Assignment 2
- Can revise Assignment 01 & turn in morning after Class 2

Between Class 2 & Class 3

- async material for Week 2
- Readings for Week 2
- Turn in Assignment 02

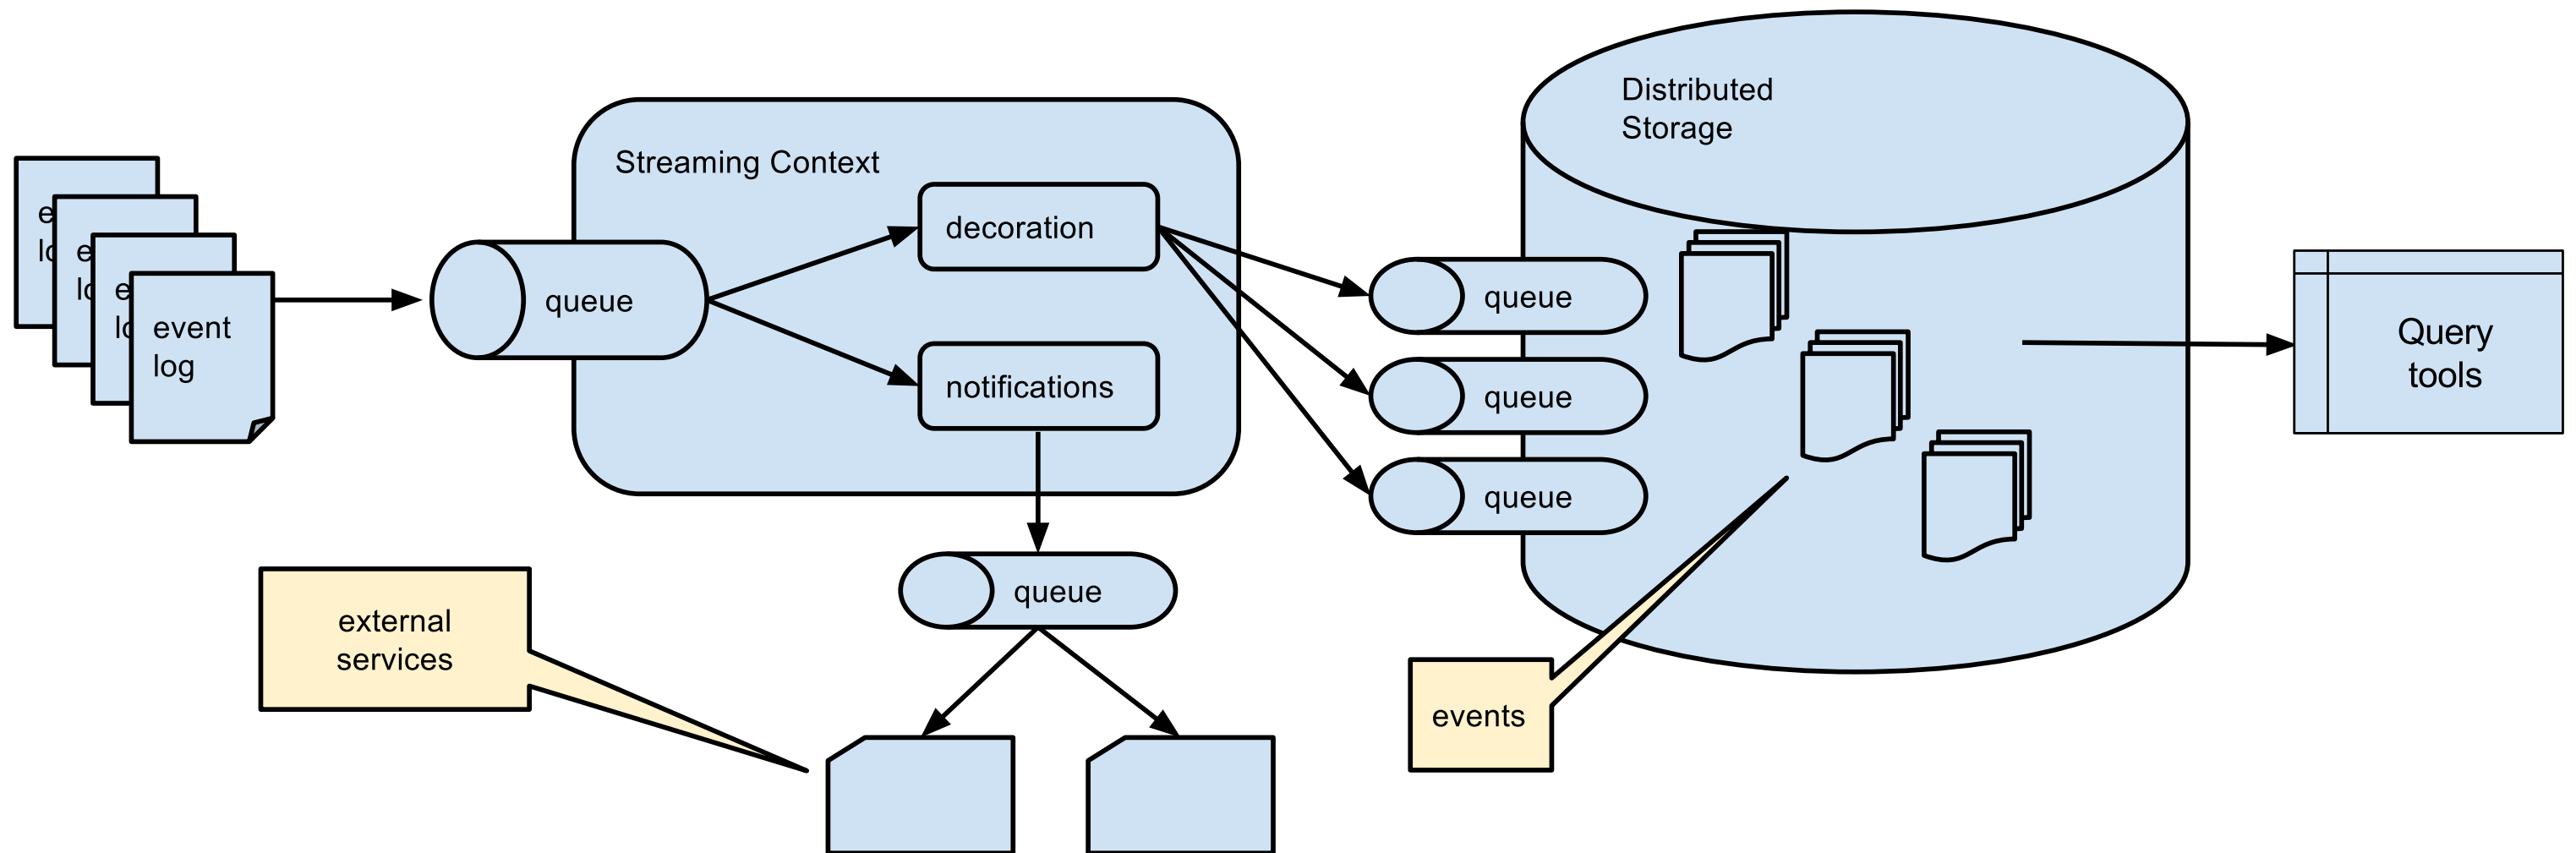
Class 3

- Groups share Assignment 02
- Preview, discussion, do google cloud platform setup and sql statements for Assignment 03
- Can revise Assignment 02 & turn in morning after Class 3

Student Projects

Student Projects

1. Querying Data
2. Tracking User Activity
3. Understanding User Behavior



Querying Data

- Use existing tools/pipeline/dataset
- Answer basic business questions

Tracking User Activity

- Use provided pipeline components
- Transform/store data
- Answer business questions
- Bonus:
 - Trigger notifications

Understanding User Behavior

- Assemble an end-to-end pipeline
- Ingest/transform/store data
- Answer comprehensive business questions
- Bonus:
 - Manage sessionization / state

Levels of Expertise

Berkeley

SCHOOL OF
INFORMATION